

Load Cifar10 image dataset and display the first 25 images to verify that loading was successful

```
In [1]: from __future__ import absolute_import, division, print_function, unicode_literals

import tensorflow as tf

from tensorflow.keras import datasets, layers, models
import matplotlib.pyplot as plt

(train_images, train_labels), (test_images, test_labels) = datasets.cifar10.load_data()

# Normalize pixel values to be between 0 and 1
train_images, test_images = train_images / 255.0, test_images / 255.0

class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
               'dog', 'frog', 'horse', 'ship', 'truck']

plt.figure(figsize=(10,10))
for i in range(25):
    plt.subplot(5,5,i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(train_images[i], cmap=plt.cm.binary)
    # The CIFAR labels happen to be arrays,
    # which is why you need the extra index
    plt.xlabel(class_names[train_labels[i][0]])
plt.show()
```

<Figure size 1000x1000 with 25 Axes>

Create our Convolutional Neural Network. Note that in Tensorflow, we create our network as a static graph and then run our data through it. We cannot change the structure of the network dynamically and the structure of the network does not depend on the data distribution.

```
In [2]: model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))
```

View our Network structure and the number of trainable parameters on each layer.

```
In [3]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 30, 30, 32)	896
-----		
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
-----		
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
-----		
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
-----		
conv2d_2 (Conv2D)	(None, 4, 4, 64)	36928
-----		
flatten (Flatten)	(None, 1024)	0
-----		
dense (Dense)	(None, 64)	65600
-----		
dense_1 (Dense)	(None, 10)	650
=====		
Total params: 122,570		
Trainable params: 122,570		
Non-trainable params: 0		
-----		

Run the network with our data.

```
In [4]: model.compile(optimizer='adam',
                    loss='sparse_categorical_crossentropy',
                    metrics=['accuracy'])

history = model.fit(train_images, train_labels, epochs=10,
                    validation_data=(test_images, test_labels))
```

Train on 50000 samples, validate on 10000 samples

Epoch 1/10

50000/50000 [=====] - 17s 349us/sample - loss: 1.5157 - accuracy: 0.4481 - val\_loss: 1.2657 - val\_accuracy: 0.5394

Epoch 2/10

50000/50000 [=====] - 12s 246us/sample - loss: 1.1629 - accuracy: 0.5875 - val\_loss: 1.0821 - val\_accuracy: 0.6147

Epoch 3/10

50000/50000 [=====] - 12s 245us/sample - loss: 1.0172 - accuracy: 0.6408 - val\_loss: 0.9818 - val\_accuracy: 0.6596

Epoch 4/10

50000/50000 [=====] - 12s 238us/sample - loss: 0.9169 - accuracy: 0.6760 - val\_loss: 0.9709 - val\_accuracy: 0.6640

Epoch 5/10

50000/50000 [=====] - 12s 239us/sample - loss: 0.8329 - accuracy: 0.7081 - val\_loss: 0.9215 - val\_accuracy: 0.6820

Epoch 6/10

50000/50000 [=====] - 12s 240us/sample - loss: 0.7768 - accuracy: 0.7251 - val\_loss: 0.8841 - val\_accuracy: 0.6950

Epoch 7/10

50000/50000 [=====] - 14s 285us/sample - loss: 0.7205 - accuracy: 0.7455 - val\_loss: 0.8969 - val\_accuracy: 0.6931

Epoch 8/10

50000/50000 [=====] - 13s 251us/sample - loss: 0.6771 - accuracy: 0.7626 - val\_loss: 0.8878 - val\_accuracy: 0.6987

Epoch 9/10

50000/50000 [=====] - 12s 243us/sample - loss: 0.6390 - accuracy: 0.7746 - val\_loss: 0.9177 - val\_accuracy: 0.6962

Epoch 10/10


50000/50000 [=====] - 12s 245us/sample - loss: 0.6017 - accuracy: 0.7866 - val\_loss: 0.8685 - val\_accuracy: 0.7109

Plot the train and validation accuracy

```
In [6]: plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0.5, 1])
plt.legend(loc='lower right')

test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
```

10000/1 - 1s - loss: 0.8628 - accuracy: 0.7109



Print the final test accuracy

```
In [7]: print(test_acc)
```

0.7109

```
In [8]: #Create the confusion matrix
```

```
File "<ipython-input-8-ec6ee38ea864>", line 1
    Create the confusion matrix
          ^
SyntaxError: invalid syntax
```

```
In [9]: from sklearn.metrics import classification_report, confusion_matrix
import numpy as np

pred=model.predict(train_images)
y_pred = np.argmax(pred, axis=1)
print('Confusion Matrix')
print(confusion_matrix(train_labels, y_pred))
#target_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
#                'dog', 'frog', 'horse', 'ship', 'truck']
```

Confusion Matrix

```
[[4161  85 260  41  54  9 12 28 270 80]
 [ 32 4707 11 10 12 8 13 3 122 82]
 [ 162 24 3999 138 259 103 129 116 57 13]
 [ 59 34 360 3314 241 488 215 162 83 44]
 [ 82 12 342 177 3899 72 71 309 25 11]
 [ 20 23 290 870 168 3195 110 267 33 24]
 [ 19 29 241 168 188 48 4235 20 36 16]
 [ 29 15 128 99 99 88 17 4481 17 27]
 [ 135 42 35 24 6 6 12 7 4707 26]
 [ 96 461 42 28 7 7 17 32 121 4189]]
```

```
In [ ]: #Class 1 mistaken for class 3 the most
#Class 2 mistaken for class 9 the most
#Class 3 mistaken for class 5 the most -> Implies Class 3's features aren't di
stinguishable enough
#Class 4 mistken for class 6 the most
#Class 5 mistaken for class 3 the most
#Class 6 mistaken for class 4 the most
#Class 7 mistaken for class 3 the most
#Class 8 mistaken for class 3 the most
#Class 9 mistaken for class 1 the most
#Class 10 mistaken for class 2 the most
```